

Midterm Paper 1 Answers

(WhatsAppImage2025-03-11at10.35.07_753228)

This document provides detailed answers to the questions from the first midterm paper image provided.

Question 1: Differentiate Synchronous and asynchronous executions with diagrams. (10 Marks)

Concept: Synchronous vs. Asynchronous Execution Models in Distributed Systems.

Syllabus Relevance: Module 1 (Synchronous versus asynchronous executions, Design issues and challenges).

Answer:

In distributed systems, the distinction between synchronous and asynchronous execution models is fundamental, primarily defined by assumptions about process execution speeds, message delivery times, and clock drift rates. These assumptions significantly impact the design and complexity of distributed algorithms.

Synchronous Distributed System: A synchronous distributed system is characterized by known, finite upper bounds on process execution speeds, message transmission delays, and the rate at which local clocks drift from real-time. Specifically:

1. **Bounded Message Delay:** There is a known maximum time t_{max} for any message sent between two connected processes to be received.
2. **Bounded Execution Steps:** The time taken for a process to execute a single step or instruction has a known upper bound.
3. **Bounded Clock Drift:** The rate at which any process's local clock deviates from a perfect reference clock is bounded by a known constant.

These bounds allow processes to operate in lock-step rounds. Processes can wait for a certain duration, knowing that any message sent in the previous round will have arrived. This simplifies algorithm design, particularly for tasks requiring coordination or agreement, as timeouts can reliably detect failures (a process not responding within the maximum time is considered failed). Consensus, for example, is solvable deterministically in synchronous systems even with certain types of failures.

Diagrammatic Representation (Conceptual): Imagine processes P1, P2, P3 executing in rounds. In each round r , a process performs computation, sends messages, and then

waits for a duration Δt (calculated based on bounds) before starting round $r+1$. All messages sent in round r are guaranteed to arrive before the end of the waiting period Δt , allowing processes to start round $r+1$ with complete information from round r .

Time ->

P1: [Round 1 Comp/Send] --Wait Δt --> [Round 2 Comp/Send] --Wait Δt --> ...

P2: [Round 1 Comp/Send] --Wait Δt --> [Round 2 Comp/Send] --Wait Δt --> ...

P3: [Round 1 Comp/Send] --Wait Δt --> [Round 2 Comp/Send] --Wait Δt --> ...

| <----- Round 1 -----> | <----- Round 2 -----> |

(Messages sent in Round 1 arrive before Round 2 starts)

Asynchronous Distributed System: An asynchronous distributed system makes no assumptions about the relative timing of events. Specifically: 1. **Unbounded Message Delay:** Message transmission delays can be arbitrarily long, though finite (messages are eventually delivered if the recipient is correct and the network doesn't partition permanently). 2. **Unbounded Execution Steps:** The time taken for a process to execute a step is unpredictable. 3. **Unbounded Clock Drift:** Local clocks can drift at arbitrary rates relative to each other and real-time.

In this model, it is impossible to use timeouts to reliably detect process failures, as a delayed message is indistinguishable from a message sent by a crashed process. Algorithms must rely solely on message arrivals and cannot assume lock-step execution. This makes designing algorithms for problems like consensus significantly harder; for instance, the FLP impossibility result shows that deterministic consensus is impossible in a purely asynchronous system with even one potential crash failure. Algorithms often rely on message ordering guarantees (like FIFO or causal order) or failure detectors to achieve progress.

Diagrammatic Representation (Conceptual): Processes P1, P2, P3 execute independently. Messages (m_1 , m_2) sent can experience vastly different delays. P1 might send m_1 to P2, and P2 might send m_2 to P3 much later, with m_1 arriving after m_2 is sent, or vice-versa. There are no synchronized rounds.

Time ->

P1: Event -> Send(m_1) -> Event -> ...

|
P2: Event -> ... -> Receive(m_1) -> Send(m_2) -> ...

|
P3: Event -> ... -> Receive(m_2) -> ...

(Message delays are arbitrary; no fixed rounds)

Key Differences Summarized:

Feature	Synchronous System	Asynchronous System
Message Delay	Bounded	Unbounded
Process Speed	Bounded step time	Unbounded step time
Clock Drift	Bounded	Unbounded
Failure Detection	Possible via timeouts	Impossible via timeouts
Execution Model	Lock-step rounds possible	Independent, event-driven
Algorithm Design	Simpler (e.g., consensus)	More complex (e.g., FLP)

Question 2: In the above space-time diagram of a distributed execution, global state G1 consisting of the given local states. Does the global state G1 is consistent? Justify your answer. $G1 = \{ L1(e1^3), L2(e2^2), L3(e3^1) \}$ (10 Marks)

Concept: Consistent Global State (Consistent Cut), Happened-Before Relation.

Syllabus Relevance: Module 1 (Global state, Cuts), Module 2 (Snapshot algorithms).

Answer:

A global state, represented as a cut across the space-time diagram of a distributed execution, is considered **consistent** if, for every message m delivered (received) by a process included in the cut, the corresponding send event of m by the source process is also included in the cut. In simpler terms, a cut is consistent if no message crosses the cut line from the future (events after the cut) to the past (events before the cut).

The given global state is $G1 = \{ L1(e1^3), L2(e2^2), L3(e3^1) \}$. This corresponds to a cut $C1$ that passes through the state intervals immediately following events $e1^3$ on process $P1$, $e2^2$ on process $P2$, and $e3^1$ on process $P3$.

Let's analyze the messages exchanged relative to this cut $C1$:

1. **Message m_{12} (from $P1$ to $P2$):** Sent by event $e1^1$ (before $e1^3$ on $P1$, so included in the past of $C1$). Received by event $e2^2$ (the last event included on $P2$ by $C1$). Since both send and receive are before or at the cut, this message does not violate consistency.

2. **Message m13 (from P1 to P3):** Sent by event $e1^2$ (before $e1^3$ on P1, so included in the past of C1). Received by event $e3^2$ (after $e3^1$ on P3, so not included in the past of C1). This message crosses the cut from past to future, which is allowed and does not violate consistency.
3. **Message m21 (from P2 to P1):** Sent by event $e2^1$ (before $e2^2$ on P2, so included in the past of C1). Received by event $e1^3$ (the last event included on P1 by C1). Since both send and receive are before or at the cut, this message does not violate consistency.
4. **Message m31 (from P3 to P1):** Sent by event $e3^1$ (the last event included on P3 by C1). Received by event $e1^4$ (after $e1^3$ on P1, so not included in the past of C1). This message crosses the cut from past to future, which is allowed and does not violate consistency.
5. **Message m32 (from P3 to P2):** Sent by event $e3^1$ (the last event included on P3 by C1). Received by event $e2^3$ (after $e2^2$ on P2, so not included in the past of C1). This message crosses the cut from past to future, which is allowed and does not violate consistency.

Crucially, let's re-examine the diagram carefully for any message received before the cut whose send event happened after the cut. Looking at the cut defined by $\{e1^3, e2^2, e3^1\}$: * P1 receives m21 at $e1^3$. m21 was sent at $e2^1$ (before $e2^2$). OK. * P2 receives m12 at $e2^2$. m12 was sent at $e1^1$ (before $e1^3$). OK. * P3 is cut after $e3^1$. It has not received any messages up to this point.

All messages received up to the cut points $\{e1^3, e2^2, e3^1\}$ were sent before the corresponding cut points on their respective sender processes. No message crosses the cut from the future to the past.

Justification: The global state $G1 = \{ L1(e1^3), L2(e2^2), L3(e3^1) \}$ represents a consistent cut of the distributed execution. This is because for every event e included in the state (i.e., e happened-before or is the cut event itself), if e is a receive event $Receive(m)$, then the corresponding send event $Send(m)$ is also included in the state (i.e., $Send(m)$ happened-before the cut event on the sender process). We have verified this for all messages received up to the cut points: m12 (sent $e1^1$, rcv $e2^2$) and m21 (sent $e2^1$, rcv $e1^3$). No message reception included in $G1$ has its corresponding send event excluded from $G1$.

Therefore, the global state G1 is consistent.

Question 3: Consider the following three process A-execution diagram and determine whether it obeys causal order or not with justification. (10 Marks)

Concept: Causal Order (CO) in Message Delivery.

Syllabus Relevance: Module 2 (Message ordering paradigms, Causal order (CO)).

Answer:

Causal Order (CO) is a message ordering paradigm in distributed systems, particularly relevant for group communication or multicast messages. It states that if the sending of message m_1 happened-before the sending of message m_2 , then any process p that delivers both messages must deliver m_1 before it delivers m_2 . The happened-before relation (\rightarrow) is defined by Lamport: $a \rightarrow b$ if (i) a and b are events in the same process and a occurs before b , or (ii) a is the sending of a message m and b is the reception of m , or (iii) there exists an event c such that $a \rightarrow c$ and $c \rightarrow b$ (transitivity).

Let's analyze the provided A-execution diagram (which seems to depict asynchronous point-to-point communication rather than multicast, but the principle of causality still applies to the order of events observed by processes):

- **Process P1:** Sends m_1 (event s_1), then later sends m_3 (event s_3).
- **Process P2:** Sends m_2 (event s_2).
- **Process P3:** Receives m_1 (event r_1), then receives m_2 (event r_2), then receives m_3 (event r_3).

We need to check if any causal relationship between send events is violated by the receive order at P3.

1. **Causality between m_1 and m_3 :** Event s_1 (send m_1) happens before s_3 (send m_3) on process P1. Thus, $s_1 \rightarrow s_3$.
2. **Delivery Order at P3:** P3 delivers m_1 (at r_1) and then delivers m_3 (at r_3). Since $s_1 \rightarrow s_3$, CO requires that m_1 be delivered before m_3 at any common destination. P3 delivers m_1 before m_3 . This order respects causality.

3. Other Potential Causal Chains:

- Is there a causal path from s_2 (send m_2) to s_1 (send m_1)? No direct path.
- Is there a causal path from s_2 (send m_2) to s_3 (send m_3)? No direct path.
- Is there a causal path from s_1 (send m_1) to s_2 (send m_2)? No direct path.
- Is there a causal path from s_3 (send m_3) to s_2 (send m_2)? No direct path.

Let's consider the full happened-before relation based on the diagram: * $s_1 \rightarrow s_3$ (on P1)
* $s_1 \rightarrow r_1$ (message m1) * $s_2 \rightarrow r_2$ (message m2) * $s_3 \rightarrow r_3$ (message m3) * $r_1 \rightarrow r_2 \rightarrow r_3$ (on P3)

From these, we can infer causal chains: * $s_1 \rightarrow r_1 \rightarrow r_2$ * $s_1 \rightarrow r_1 \rightarrow r_3$ * $s_1 \rightarrow s_3 \rightarrow r_3$ * $s_2 \rightarrow r_2 \rightarrow r_3$

Now, let's re-evaluate the CO condition specifically for messages delivered at P3: *

Messages m1 and m3: We established $s_1 \rightarrow s_3$. P3 delivers m1 (at r1) before m3 (at r3).

This respects CO. * **Messages m1 and m2:** There is no causal relationship ($s_1 \rightarrow s_2$ or $s_2 \rightarrow s_1$) shown between the send events. Therefore, CO imposes no constraint on their delivery order relative to each other. P3 delivers m1 (r1) then m2 (r2).

* **Messages m2 and m3:** There is no causal relationship ($s_2 \rightarrow s_3$ or $s_3 \rightarrow s_2$) shown between the send events. Therefore, CO imposes no constraint on their delivery order relative to each other. P3 delivers m2 (r2) then m3 (r3).

Justification: The execution obeys Causal Order. The only pair of messages with a causal dependency between their send events is (m1, m3), where $\text{send}(m_1)$ happened-before $\text{send}(m_3)$. Process P3 delivers m1 before delivering m3, satisfying the requirement of Causal Order. For all other pairs of messages delivered at P3 (m1, m2) and (m2, m3), there is no causal relationship between their respective send events, so any delivery order is permissible under CO.

Therefore, the execution shown obeys Causal Order.

Question 4: Consider the A-execution diagram in question 3. Do you think message ordering is satisfied. Justify your answer. (10 Marks)

Concept: Message Ordering Paradigms (FIFO, Causal, Total).

Syllabus Relevance: Module 2 (Message ordering paradigms).

Answer:

This question asks if

"message ordering" is satisfied in the execution shown in Question 3. Message ordering isn't a single property but rather a set of different guarantees, such as FIFO (First-In, First-Out), Causal Order (CO), and Total Order (TO).

Let's evaluate the common ordering guarantees based on the diagram:

1. **FIFO Order:** This guarantees that messages sent between the same pair of processes (sender and receiver) are delivered in the order they were sent.
 - Consider the channel P1 → P3: P1 sends m1 (event s1) and later sends m3 (event s3). Process P3 receives m1 (event r1) and later receives m3 (event r3). The delivery order (m1 then m3) matches the sending order (m1 then m3). So, FIFO order is satisfied for the P1 → P3 channel.
 - Consider the channel P2 → P3: P2 sends only one message, m2. FIFO is trivially satisfied.
 - No other channels have multiple messages in this diagram.
 - Therefore, FIFO order is satisfied for all relevant communication channels shown.
2. **Causal Order (CO):** As established in the answer to Question 3, if $\text{send}(m1) \rightarrow \text{send}(m2)$, then any process delivering both must deliver m1 before m2.
 - We found that $s1 \rightarrow s3$. P3 delivers m1 (at r1) before m3 (at r3). This respects CO.
 - There were no other causal dependencies between send events for messages delivered at P3.
 - Therefore, Causal Order is satisfied.
3. **Total Order (TO):** This guarantees that all processes that deliver a set of messages deliver them in the same relative order. This is typically applied to multicast messages. The diagram shows point-to-point messages, and only P3 receives messages m1, m2, and m3. P3 delivers them in the sequence m1, m2, m3. Without other recipients, we cannot fully verify or falsify Total Order across multiple processes. However, there is no violation shown at the single recipient P3.

Justification: Message ordering, interpreted as either FIFO or Causal Order, is satisfied in the given A-execution diagram. * **FIFO:** Messages sent from P1 to P3 (m1, then m3) are received in the same order (r1, then r3). * **Causal:** The only causal dependency between send events ($\text{send}(m1) \rightarrow \text{send}(m3)$) is respected by the delivery order at P3 ($\text{deliver}(m1)$ before $\text{deliver}(m3)$).

Since the fundamental ordering properties (FIFO and Causal) are met based on the diagram, we can conclude that message ordering is satisfied.

Question 5: Explain the three basic approaches for implementing distributed mutual exclusion with system model. (10 Marks)

Concept: Distributed Mutual Exclusion (DME) Approaches, System Model.

Syllabus Relevance: Module 3 (Distributed mutual exclusion algorithms, System model).

Answer:

Distributed Mutual Exclusion (DME) ensures that at most one process can access a shared resource (Critical Section - CS) at any given time in a distributed system where processes communicate only via message passing.

System Model: We typically assume a distributed system consisting of N autonomous processes (P_1, P_2, \dots, P_N) connected by a communication network. Key assumptions often include: * **Processes:** Processes operate asynchronously at potentially different speeds. They can fail (e.g., crash failures). * **Communication:** Processes communicate solely by exchanging messages. The network might be assumed to be reliable (messages are eventually delivered without corruption) or unreliable. Message delivery might be FIFO ordered between pairs of processes or unordered. The system might be synchronous (bounded message delays) or asynchronous (unbounded delays). * **No Shared Memory:** Processes do not share physical memory.

Basic Approaches for DME:

1. Non-Token-Based (Permission-Based) Approach:

- **Idea:** A process wishing to enter the CS must request permission from some set of other processes and wait until sufficient permissions are granted. Timestamps (like Lamport's logical clocks or vector clocks) or sequence numbers are typically used to order requests and resolve conflicts fairly.
- **Mechanism:** When P_i wants to enter the CS, it sends a REQUEST message (often timestamped) to a set of other processes (e.g., all other processes in Lamport's or Ricart-Agrawala's algorithms). A receiving process P_j grants permission (sends a REPLY) based on its own state (whether it's in the CS, requesting the CS, or idle) and the priority of the requests (e.g., lower timestamp has higher priority). P_i enters the CS only after receiving REPLY messages from the required set of processes. Upon exiting the CS, P_i sends RELEASE messages or informs deferred requesters.
- **Examples:** Lamport's Algorithm, Ricart-Agrawala Algorithm.

- **Diagram (Conceptual - Ricart-Agrawala):** `` P1 (Wants CS) -- REQUEST(ts1)--> P2 P1 (Wants CS) --REQUEST(ts1)--> P3

P2 (Idle) --REPLY--> P1 P3 (Wants CS, ts3 > ts1) --REPLY--> P1

P1 (Receives all Replies) --> Enter CS --> Exit CS --> Send deferred Replies ``

2. Token-Based Approach:

- **Idea:** A unique token exists in the system. Only the process currently holding the token is allowed to enter the CS. Mutual exclusion is guaranteed because only one process can possess the token at a time.
- **Mechanism:** Processes are typically organized in a logical structure (e.g., a ring or a tree). A process wanting to enter the CS requests the token. If it doesn't have the token, the request is forwarded along the logical structure towards the current token holder. When the token holder exits the CS (or if it wasn't using it), it passes the token to the next requesting process according to the algorithm's rules (e.g., along the ring, down the tree, or directly in broadcast-based algorithms).
- **Examples:** Suzuki-Kasami's Broadcast Algorithm, Raymond's Tree-Based Algorithm.
- **Diagram (Conceptual - Ring):** `` P1 --> P2 --> P3 --> P4 --> P1 (Logical Ring) (Token initially at P1)

P3 (Wants CS) --> Sends Request towards P1 P1 (Exits CS/Idle) --> Sends Token --> P2 --> P3 P3 (Receives Token) --> Enter CS --> Exit CS --> Sends Token to next requester/neighbor ``

3. Quorum-Based Approach:

- **Idea:** A process wishing to enter the CS requests permission from only a subset of processes, called a quorum (or request set). These quorums are carefully constructed such that any two quorums in the system have at least one process in common (non-empty intersection property).
- **Mechanism:** When P_i wants to enter the CS, it sends REQUEST messages only to the processes in its quorum (R_i). A process P_j in R_i grants permission (sends REPLY) only if it hasn't granted permission to another process already (it locks itself for P_i). P_i enters the CS only after receiving permission from all processes in its quorum R_i . The intersection property ensures that no two processes can get permission from their respective quorums simultaneously, because the common member(s) will grant permission to only one of them at

a time. Upon exiting, P_i sends RELEASE messages to its quorum members, allowing them to grant permission to others.

- **Example:** Maekawa's Algorithm (uses finite projective planes to construct quorums of size approx. \sqrt{N}).
- **Diagram (Conceptual - Maekawa):** `` ` N=7 processes. Quorums R_i, R_j (size $\sim \sqrt{7} \sim 3$). Intersection: $R_i \cap R_j \neq \emptyset$ (e.g., P_k is in both)

P_1 (Wants CS) --REQUEST--> Quorum $R_1 = \{P_1, P_2, P_4\}$ P_5 (Wants CS) --REQUEST--> Quorum $R_5 = \{P_2, P_5, P_6\}$

P_1 gets REPLY from P_1, P_4 . P_5 gets REPLY from P_5, P_6 .

P_2 (Common member) receives requests from P_1 and P_5 . P_2 grants permission (REPLY) to only one (e.g., P_1 based on timestamp). P_2 defers REPLY to P_5 .

P_1 receives all replies --> Enter CS. P_5 waits for reply from P_2 .

P_1 (Exits CS) --RELEASE--> $R_1 = \{P_1, P_2, P_4\}$ P_2 (Receives RELEASE from P_1) --> Sends REPLY to P_5 . P_5 (Receives reply from P_2) --> Enter CS. `` `

Each approach has different trade-offs regarding message complexity, synchronization delay, fault tolerance, and ease of implementation.